

Differentiable Parameter Recovery for ECCO-Darwin Biogeochemistry: A BINN-Style Framework and a Structural 5/6 Ceiling

Internal white paper — DarwinDiff project

May 22, 2026

Abstract

ECCO-Darwin’s six-parameter biogeochemistry calibration uses Green’s functions: expensive perturb-and-fit Jacobians that don’t transfer when new observations come online. We replace it with a Biogeochemistry-Informed Neural Network (BINN) style framework adapted to marine biogeochemistry. A small per-cell neural network maps environmental covariates to the six Carroll-N parameters, sigmoid-bounded into Carroll’s published ranges and fed through a differentiable PyTorch reimplementation of Darwin’s source-minus-sink (SMS) kernel; the observation loss backpropagates end-to-end. We report the v3.1 sweep: **856 seeds in 86 single-lever configurations**. Two findings. First, the iron pair (`alpfe`, `scav_rat`) recovers reproducibly at **38/40 (95%) Cal-grade** with our best configuration. Second, a structural **5/6 ceiling**: 0/856 seeds reach all six Carroll-N parameters jointly, with only two unreproduced single-seed 5/6 events; a composition test of the two complementary 5/6 lever families regresses below either parent. The ceiling is consistent with parameter conservation under the available observations and predicts which new observation channels would unblock 6/6. Source code and full sweep outputs accompany this paper.

1 Background

ECCO-Darwin (Carroll et al., 2020, 2022) couples MITgcm physics with the Darwin biogeochemistry module on the ECCO LLC270 grid (50 vertical levels, 1992–2017+). Its biogeochemistry exposes 103 independent tunable scalar parameters in source; six are calibrated by Green’s functions (Menemenlis et al., 2005): iron dust solubility `alpfe`, iron scavenging rate `scav_rat`, small phytoplankton growth `Smallgrow`, large phytoplankton growth `Biggrow`, diatom palatability `diatomgraz`, and the PIC/POC ratio `R_PICPOC`. Carroll 2022 inherits Carroll 2020’s six calibrated values bit-for-bit.

Green’s-functions calibration is limited in two ways. Each tuned parameter requires a full perturbation-and-rerun campaign, so the calibrated subset stays small. And the calibration is implicitly tied to its observational input set — adding a new observation channel (PACE ocean color, expanding GEOTRACES) demands a fresh Jacobian campaign, not a reuse of the existing one.

The methodological precedent for our approach is BINN, the Biogeochemistry-Informed Neural Network of Xu et al. (2025): a small MLP predicts soil-carbon parameters per location, sigmoid-bounded into prior ranges, fed into a differentiable PyTorch reimplementation of CLM5 with an SOC-profile loss that backpropagates end-to-end. We adapt the pattern from terrestrial soil carbon




Figure 1: **DarwinDiff training pipeline.** A per-cell neural network (DINN) maps environmental covariates to six raw values; sigmoid bounding maps them into Carroll’s published parameter ranges; a differentiable PyTorch reimplement of Darwin’s source-minus-sink (SMS) kernel advances five to seven prognostic tracers via forward-Euler integration; the z-scored MSE loss against ECCO-Darwin v05 outputs backpropagates through the integration to the network weights. New observation channels enter as additional loss terms.

to marine biogeochemistry. Differentiable-physics neighbors that do not address parameter recovery within ocean BGC include Neural-BGC (Ouala and Lachkar, 2026) (pure-NN emulation of O_2 and NO_3) and differentiable Veros (Meunier et al., 2025).

2 Method

The full data flow is shown in Figure 1. Three components:

DINN. A per-cell 1×1 -convolutional MLP. The baseline DINN (~ 454 weights, SST input only) is used for the structural-ceiling experiments where the cleanest comparison against the global-scalar Green’s-functions class matters. The production DINNDeep (~ 9.4 k weights, SST + MLD + windspeed + latitude) is used for higher-quality within-region fits. The 1×1 kernel commits to per-cell parameter independence; spatial coupling would conflate per-cell parameter learning with smoothing.

Sigmoid bounding. Raw NN outputs are mapped to Carroll’s published parameter ranges via $\theta_i = \ell_i + (h_i - \ell_i) \sigma(\hat{\theta}_i)$ so recovered parameters stay biologically sensible regardless of NN output magnitude.

Differentiable SMS kernel. A PyTorch reimplement of Darwin’s source-minus-sink rate equations (canonically the `darwin_plankton.F` subroutine in both Carroll 2020’s Darwin 1 and the upstream `darwinproject/darwin3` repository). Five prognostic tracers in the baseline (DFe, P_s , P_l , POC, PIC); seven with the carbonate extension (adds DIC, ALK, with air-sea CO_2 flux as a diagnostic via the Follows-2006 carbonate solver (Follows et al., 2006) and the Wanninkhof-2014 gas-transfer law (Wanninkhof, 2014)). Integrated by forward-Euler ($\Delta t = 0.25$ d, 200 steps) to near-steady state. The remaining ~ 94 background parameters Carroll left untuned are held fixed at his literature values.

Training. 1500 epochs, Adam at 5×10^{-3} , z-scored MSE on the spatial pattern of each tracer field over ocean cells, autograd through the 200 integration steps. One fit completes in ~ 7 –8 min on an RTX 5090 Laptop.

Multi-area-of-interest joint training (v3.1). Three areas of interest (Equatorial Pacific, North Atlantic Subpolar, Southern Ocean Pacific) contribute joint loss; the optimizer must find Carroll-N values satisfying all three. Two architecture variants: a shared DINN with an AOI-identity input channel, and a per-AOI DINN with a soft cross-AOI consistency penalty $\lambda_{\text{consistency}}$.

Observation channels. Per-AOI chlorophyll across five phytoplankton functional types, POC, DIC, ALK, air-sea CO_2 flux; GEOTRACES IDP2025 surface and subsurface dissolved iron, GEOTRACES subsurface POC, and GEOTRACES surface biogenic silica (via a steady-state diagnostic). Loss weights are configuration-level hyperparameters varied systematically across the sweep.

3 Results

3.1 Sweep summary

Five chained waves of single-lever sweeps plus a Wave 6 composition test. **856 seeds across 86 configurations** after aggregation; Table 1 reports the headline counts.

Table 1: v3.1 sweep summary. “Cal-grade” means a recovered Carroll-N parameter falls within $\pm 40\%$ of Carroll 2020’s published optimum.

Metric	Value
Single-lever configurations	86
Total seeds	856
Seeds at 6/6 Cal-grade	0
Seeds at 5/6 Cal-grade	2 (0.23%)
Configurations preserving the iron pair at $\geq 9/10$ seeds	40
F2 Basin C iron-pair Cal-grade rate, $n=40$	38/40 (95%)

3.2 Basin C iron-pair recovery

The F2 configuration — `POSI_W=1.0`, `AOI_W_NATLSUBPOLAR=2.0`, `AOI_W_SOUTHERNOCEANPAC=2.0`, `CHL1_W_EXTRA=3.0`, otherwise the Basin C 3-AOI base — recovers the iron pair (`alpfe`, `scav_rat`) at **38/40 (95%)** Cal-grade across four independent 10-seed batches (10/10, 10/10, 10/10, 8/10). This is the first reproducible joint recovery of `alpfe` and `scav_rat` in our framework. Figure 2 shows the recovered pairs against Carroll’s published optimum.

3.3 The 5/6 ceiling

The headline empirical finding: **0/856 seeds reach all six Carroll-N parameters jointly at Cal-grade**. Only two seeds reach 5/6, both unreproduced. Figure 3 shows the Cal-grade count distribution; Figure 4 shows per-parameter recovery spread.

The two 5/6 events recover *complementary* parameter subsets: `w2e_peranoi_lam0.1` seed 3 recovers `alpfe` (Excellent) + `scav_rat` + `Smallgrow` + `Biggrow` + `diatomgraz` but misses




Figure 2: Iron-pair recovery at the F2 configuration across $n=40$ seeds in four independent batches. The red star marks Carroll 2020’s published optimum; shaded regions show $\pm 10\%$ (Excellent) and $\pm 40\%$ (Cal-grade) bounds.

R_PICPOC; c_chl40_posi15 seed 9 recovers `alpfe` + `scav_rat` (Excellent) + `Smallgrow` + `Biggrow` + R_PICPOC but misses `diatomgraz`. At $n=20$ each, both produce only 1/20.

3.4 Composition test fails

The natural 6/6 candidate stacks both lever families. Result at $n=10$ (Figure 5): **0/10 at 5/6**, mean Cal-grade count **2.00** — regressing below either parent (2.40 and 2.70). Iron pair survives (9/10) but R_PICPOC and `diatomgraz` both drift and the small-cell phytoplankton parameters regress. The intervention families interfere; the optimizer redistributes the residual error rather than letting both target recoveries hold.

3.5 Binary observation-anchor mutex

Any non-zero PIC absolute anchor (PIC_ABS_W, tested down to 0.02, with paired POC anchors at multiple ratios) wipes iron-pair recovery to 0/10, *regardless of magnitude*. POC anchor alone also kills the iron pair, with a different downstream basin geometry. The mutex is binary in PIC presence, not graduated. Mechanism: the PIC absolute anchor jointly constrains $R_{\text{PIC:POC}} \cdot M_{\text{tot}}$ via a tight basin, pulling the optimizer to a region where R_PICPOC is Cal-grade but the iron pair must drift to balance. This is the most actionable predictive finding: whether a candidate observation channel induces mutual mutex with iron recovery is testable a priori within DarwinDiff’s framework.




Figure 3: Cal-grade count distribution across 856 seeds. Zero seeds reach 6/6 (structural ceiling); two seeds reach 5/6, both unreproduced at $n=20$.

4 Discussion

The 5/6 ceiling is supported by four independent pieces of evidence:

1. Zero 6/6 events across 856 seeds in 86 single-lever configurations.
2. `w2e_peraoi_lam0.1` at $n=20$ produces only 1/20 at 5/6.
3. `c_chl40_posi15` at $n=20$ produces only 1/20 at 5/6.
4. The Wave 6 composition test of the two complementary 5/6 lever families fails to reach 5/6 (0/10), regressing below either parent.

The framing that fits the evidence is **parameter conservation**: the available observational constraints support approximately five effective degrees of freedom over the 6-parameter Carroll-N surface, and the sixth is always the residual sink — loss weighting decides which. This is not a method failure but a diagnostic. The framework predicts what would unblock 6/6: observation channels orthogonal to the current set (satellite PIC, additional GEOTRACES traverses, etc.) — each new candidate enters as a loss term and is testable a priori, no Jacobian campaign required.

Even with the ceiling in place, DarwinDiff already buys substantial advantages over Green’s functions: a single 7-minute GPU training run replaces 12 model runs per Gauss-Newton iteration; new observation channels are a configuration change rather than a fresh perturbation campaign; per-cell parameter variation strictly contains the global-scalar class as the constant-prediction limit.

Limitations. The differentiable box model is a 5–7-tracer proxy of the 39-tracer Darwin SMS kernel; the 5-PFT $\rightarrow$ 2-PFT collapse contributes meaningfully to recovery bias against Carroll’s published values, and a regional differentiable port of the full `darwin_plankton.F` is the natural next iteration. All current fits use the 23-year time-mean (climatology) rather than time-resolved Darwin output; time-resolved fitting is cluster-gated by GPU memory.




Figure 4: Per-parameter recovery distributions across 856 seeds, normalized by Carroll 2020’s optima. The iron pair clusters near unity; the other four spread far broader.

Next steps. (i) Differentiable regional full-Darwin to eliminate the proxy gap; (ii) time-resolved fitting on cluster compute (MIT ORCD AICR / NVIDIA B200 beta); (iii) satellite observation channels (PACE ocean-color, MODIS PIC) deferred for the leapfrog companion paper to keep the v05 framing canonical here.

Code and data

Source code, sweep orchestration, per-seed result JSONs, and the figure-generation pipeline are in the project repository ([src/darwindiff/](#), [docs/paper/figures/](#)). ECCO-Darwin v05 outputs are public via the ECCO Drive and NASA GHG Center. GEOTRACES IDP2025 observations are public via the GEOTRACES Data Assembly Centre.

References

- D. Carroll, D. Menemenlis, J. F. Adkins, K. W. Bowman, H. Brix, S. Dutkiewicz, I. Fenty, M. M. Gierach, C. Hill, O. Jahn, P. Landschützer, J. M. Lauderdale, J. Liu, M. Manizza, J. D. Naviaux, C. Rödenbeck, D. S. Schimel, T. Van der Stocken, and H. Zhang. The ECCO-Darwin data-assimilative global ocean biogeochemistry model: Estimates of seasonal to multidecadal surface ocean pCO₂ and air-sea CO₂ flux. *Journal of Advances in Modeling Earth Systems*, 12(10): e2019MS001888, 2020. doi: 10.1029/2019MS001888.
- D. Carroll, D. Menemenlis, S. Dutkiewicz, J. M. Lauderdale, J. F. Adkins, K. W. Bowman, H. Brix, I. Fenty, M. M. Gierach, C. Hill, O. Jahn, P. Landschützer, M. Manizza, M. R. Mazloff, C. E. Miller, D. S. Schimel, T. Van der Stocken, A. Verdy, and H. Zhang. Attribution of space-time variability in global-ocean dissolved inorganic carbon. *Global Biogeochemical Cycles*, 36(3):e2021GB007162, 2022. doi: 10.1029/2021GB007162.




Figure 5: Wave 6 composition test versus its two parent configurations ($n=10$ each). The composed configuration regresses below either parent — direct evidence for the structural 5/6 ceiling.

- M. J. Follows, T. Ito, and S. Dutkiewicz. On the solution of the carbonate chemistry system in ocean biogeochemistry models. *Ocean Modelling*, 12(3–4):290–301, 2006. doi: 10.1016/j.ocemod.2005.05.004.
- D. Menemenlis, C. Hill, A. Adcroft, J.-M. Campin, B. Cheng, B. Ciotti, I. Fukumori, P. Heimbach, C. Henze, A. Köhl, T. Lee, D. Stammer, J. Taft, and J. Zhang. NASA supercomputer improves prospects for ocean climate research. *Eos, Transactions American Geophysical Union*, 86(9):89, 95–96, 2005. doi: 10.1029/2005EO090002.
- B. Meunier, S. Ouala, H. Frezat, J. Le Sommer, and R. Fablet. Towards a fully differentiable neural ocean model with Veros. *arXiv preprint arXiv:2511.17427*, 2025. URL <https://arxiv.org/abs/2511.17427>. EurIPS 2025 workshop submission.
- Said Ouala and Zouhair Lachkar. Neural-BGC: An observation-driven emulator for hybrid physical-biogeochemical modeling. *ESSOAr preprint, submitted to Geophysical Research Letters*, 2026. doi: 10.22541/essoar.15002003/v1. URL <https://doi.org/10.22541/essoar.15002003/v1>.
- Rik Wanninkhof. Relationship between wind speed and gas exchange over the ocean revisited. *Limnology and Oceanography: Methods*, 12(6):351–362, 2014. doi: 10.4319/lom.2014.12.351.
- Xiao Xu, Feng Tao, Zhen Zhou, Yujie Zhang, William R. Wieder, Benjamin N. Sulman, Junyi Liang, William J. Riley, and Yiqi Luo. BINN: A biogeochemistry-informed neural network for soil carbon cycle modeling. *arXiv preprint arXiv:2502.00672*, 2025. doi: 10.48550/arXiv.2502.00672. URL <https://arxiv.org/abs/2502.00672>. Verified architecture details: 4-layer MLP, 21 parameter outputs, sigmoid + linear range scaling, matrix-form differentiable CLM5 in PyTorch.